

Assignment -2
Python Programming

| | |
|---------------------|-------------------|
| Assignment Date | 23 september 2022 |
| Student Name | Mathivathani.B.G |
| Student Roll Number | 314419205021 |
| Maximum Marks | 2 Marks |

Data Visualization and Pre-processing

Question-1:

1. Load the dataset

Solution:

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
sns.set_style('darkgrid')
sns.set(font_scale=1.3)
```

In [2]:

```
df=pd.read_excel("/content/Churn_Modelling.xlsx")
```

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
sns.set_style('darkgrid')
sns.set(font_scale=1.3)
```

```
In [2]: df=pd.read_excel("/content/Churn_Modelling.xlsx")
```

Question-2:

2. Perform Below Visualizations.


- Univariate Analysis
- Bi - Variate Analysis
- Multi - Variate Analysis

Solution:


```
.
#Perform Univariate Analysis
plt.figure(figsize=(8,8))
sns.kdeplot(x=df['Balance'])
```

```
In [7]: #Perform Univariate Analysis
plt.figure(figsize=(8,8))
sns.kdeplot(x=df['Balance'])
```

Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc3f3579c50>



```
#Perform Bivariate Analysis
plt.scatter(df.Age,df.Balance)
```



```
#Perform Bivariate Analysis
df.corr()
```

Out[9]:

| | CreditScore | Age | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|-----------------|-------------|-----------|-----------|-----------|---------------|-----------|----------------|-----------------|-----------|
| CreditScore | 1.000000 | -0.003965 | 0.000842 | 0.006268 | 0.012238 | -0.005458 | 0.025651 | -0.001384 | -0.027094 |
| Age | -0.003965 | 1.000000 | -0.009997 | 0.028308 | -0.030680 | -0.011721 | 0.085472 | -0.007201 | 0.285323 |
| Tenure | 0.000842 | -0.009997 | 1.000000 | -0.012254 | 0.013444 | 0.022583 | -0.028362 | 0.007784 | -0.014001 |
| Balance | 0.006268 | 0.028308 | -0.012254 | 1.000000 | -0.304180 | -0.014858 | -0.010084 | 0.012797 | 0.118533 |
| NumOfProducts | 0.012238 | -0.030680 | 0.013444 | -0.304180 | 1.000000 | 0.003183 | 0.009612 | 0.014204 | -0.047820 |
| HasCrCard | -0.005458 | -0.011721 | 0.022583 | -0.014858 | 0.003183 | 1.000000 | -0.011866 | -0.009933 | -0.007138 |
| IsActiveMember | 0.025651 | 0.085472 | -0.028362 | -0.010084 | 0.009612 | -0.011866 | 1.000000 | -0.011421 | -0.156128 |
| EstimatedSalary | -0.001384 | -0.007201 | 0.007784 | 0.012797 | 0.014204 | -0.009933 | -0.011421 | 1.000000 | 0.012097 |
| Exited | -0.027094 | 0.285323 | -0.014001 | 0.118533 | -0.047820 | -0.007138 | -0.156128 | 0.012097 | 1.000000 |

```
#Perform Bivariate Analysis
```

```
import statsmodels.api as sm
```

```
#define response variable  
y = df['CreditScore']
```

```
#define explanatory variable  
x = df[['EstimatedSalary']]
```

```
#add constant to predictor variables  
x = sm.add_constant(x)
```

```
#fit linear regression model  
model = sm.OLS(y, x).fit()
```

```
#view model summary  
print(model.summary())
```

```
=====
                        OLS Regression Results
=====
```

| | | | |
|-------------------|------------------|---------------------|-----------|
| Dep. Variable: | CreditScore | R-squared: | 0.000 |
| Model: | OLS | Adj. R-squared: | -0.000 |
| Method: | Least Squares | F-statistic: | 0.01916 |
| Date: | Thu, 29 Sep 2022 | Prob (F-statistic): | 0.890 |
| Time: | 14:58:55 | Log-Likelihood: | -59900. |
| No. Observations: | 10000 | AIC: | 1.198e+05 |
| Df Residuals: | 9998 | BIC: | 1.198e+05 |
| Df Model: | 1 | | |
| Covariance Type: | nonrobust | | |

```
=====
```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|-----------------|------------|----------|---------|-------|-----------|----------|
| ----- | ----- | ----- | ----- | ----- | ----- | ----- |
| const | 650.7617 | 1.940 | 335.407 | 0.000 | 646.958 | 654.565 |
| EstimatedSalary | -2.326e-06 | 1.68e-05 | -0.138 | 0.890 | -3.53e-05 | 3.06e-05 |

```
=====
```

| | | | |
|----------------|---------|-------------------|----------|
| Omnibus: | 132.939 | Durbin-Watson: | 2.014 |
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | 84.242 |
| Skew: | -0.072 | Prob(JB): | 5.10e-19 |
| Kurtosis: | 2.574 | Cond. No. | 2.32e+05 |

```
=====
```

Notes:

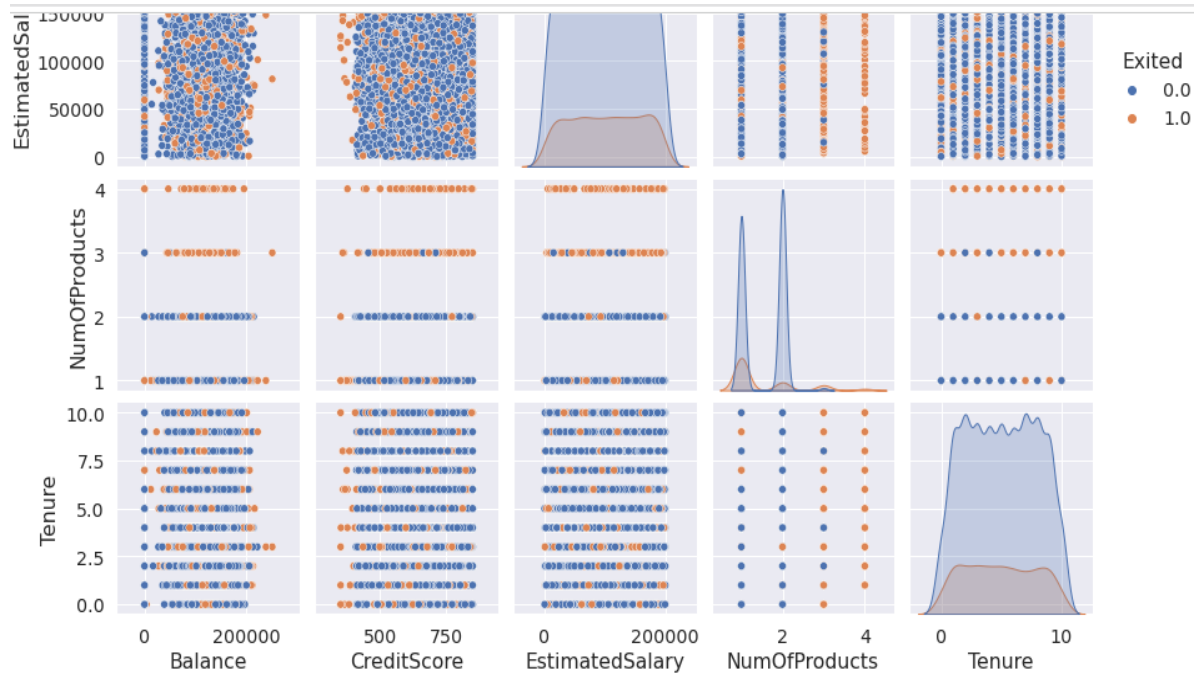
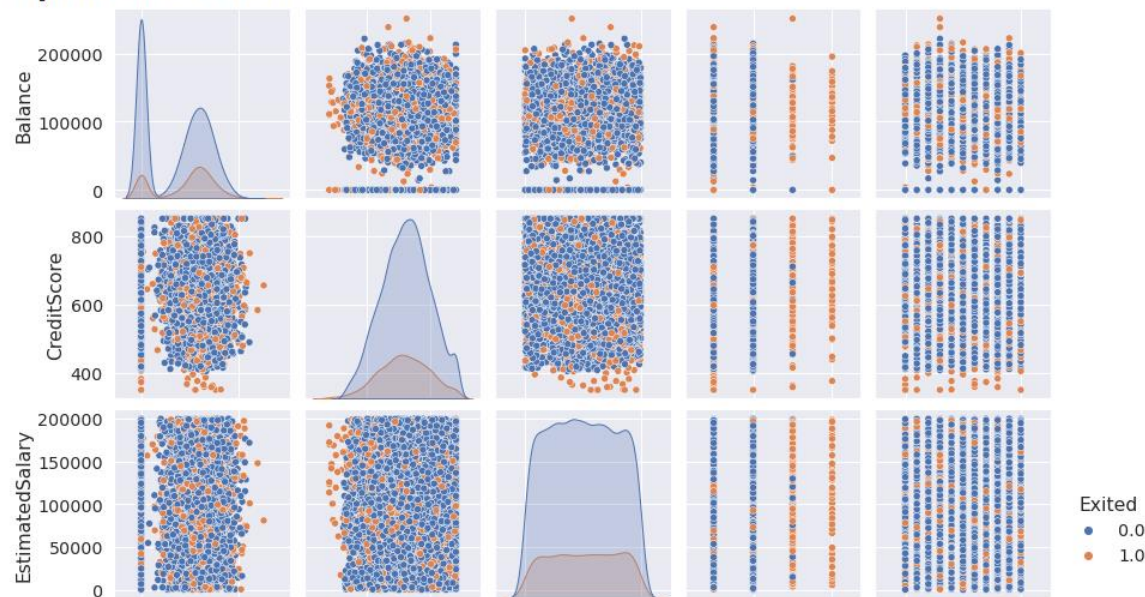
- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.32e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
#Perform Multivariate Analysis
```

```
plt.figure(figsize=(4,4))
```

```
sns.pairplot(data=df[["Balance","CreditScore","EstimatedSalary","NumOfProducts","Tenure","Exited"]],hue="Exited")
```

<Figure size 288x288 with 0 Axes>



Question-3:

3. Perform descriptive statistics on the dataset.

Solution:

```
#Perform Descriptive Statistics
df=pd.DataFrame(df)
print(df.sum())
```

```

CreditScore      6505288.0
Geography        FranceSpainFranceFranceSpainSpainFranceGermany...
Gender            FemaleFemaleFemaleFemaleFemaleMaleMaleFemaleMa...
Age              389218.0
Tenure           50128.0
Balance          764858892.88
NumOfProducts    15302.0
HasCrCard        7055.0
IsActiveMember   5151.0
EstimatedSalary  1000902398.81
Exited           2037.0
dtype: object
```

```
#Perform Descriptive Statistics
print("----Sum Value----")
print(df.sum(1))
print("-----")
print("-----Product Value----")
print(df.prod())
print("-----")
```

```

----Sum Value----
0      102015.88
1      197002.44
2      274149.37
3       94567.63
4      205492.92
...
9995    97088.64
9996   159633.38
9997    42840.58
9998   168784.83
9999   169159.57
Length: 10000, dtype: float64
-----
-----Product Value-----
CreditScore      inf
Age              inf
Tenure           0.0
Balance          0.0
NumOfProducts    inf
HasCrCard        0.0
IsActiveMember   0.0
EstimatedSalary  inf
Exited           0.0
dtype: float64
-----
```

```
#Perform Descriptive Statistics
print("-----Mean Value-----")
print(df.mean())
print("-----")
```

```

print("-----Median Value-----")
print(df.median())
print("-----")
print("-----Mode Value-----")
print(df.mode())
print("-----")
-----Mean Value-----
CreditScore      650.528800
Age              38.921800
Tenure           5.012800
Balance          76485.889288
NumOfProducts    1.530200
HasCrCard        0.705500
IsActiveMember   0.515100
EstimatedSalary  100090.239881
Exited           0.203700
dtype: float64
-----
-----Median Value-----
CreditScore      652.000
Age              37.000
Tenure           5.000
Balance          97198.540
NumOfProducts    1.000
HasCrCard        1.000
IsActiveMember   1.000
EstimatedSalary  100193.915
Exited           0.000
dtype: float64
-----
-----Mode Value-----
   CreditScore Geography Gender   Age  Tenure  Balance  NumOfProducts
0         850.0    France   Male  37.0     2.0     0.0             1.0

   HasCrCard  IsActiveMember  EstimatedSalary  Exited
0         1.0             1.0         24924.92     0.0
-----

```

Question-4:

4.Handle the Missing values

Solution:

```

#Handling with missing Values
df.isnull().values;
#Checking values are null

#Handling with missing Values
df.notnull()#Checking values are not null

```

Out[16]:

| | CreditScore | Geography | Gender | Age | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|------|-------------|-----------|--------|------|--------|---------|---------------|-----------|----------------|-----------------|--------|
| 0 | True | True | True | True | True | True | True | True | True | True | True |
| 1 | True | True | True | True | True | True | True | True | True | True | True |
| 2 | True | True | True | True | True | True | True | True | True | True | True |
| 3 | True | True | True | True | True | True | True | True | True | True | True |
| 4 | True | True | True | True | True | True | True | True | True | True | True |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9995 | True | True | True | True | True | True | True | True | True | True | True |
| 9996 | True | True | True | True | True | True | True | True | True | True | True |
| 9997 | True | True | True | True | True | True | True | True | True | True | True |
| 9998 | True | True | True | True | True | True | True | True | True | True | True |
| 9999 | True | True | True | True | True | True | True | True | True | True | True |

10000 rows × 11 columns


Question-5:

5. Find the outliers and replace the outliers

Solution:

```
#Find outliers & replace the outliers
sns.boxplot(df['Balance'])
```

Out[17]:



```
#Find outliers & replace the outliers
print(np.where(df['Balance']>100000))
(array([ 2,  4,  5, ..., 9987, 9993, 9999]),)
```

```
#Find outliers & replace the outliers
from scipy import stats
import numpy as np
```

In [19]:

```

z = np.abs(stats.zscore(df["EstimatedSalary"]))
print(z)

```

```

0      0.021886
1      0.216534
2      0.240687
3      0.108918
4      0.365276
...
9995   0.066419
9996   0.027988
9997   1.008643
9998   0.125231
9999   1.076370
Name: EstimatedSalary, Length: 10000, dtype: float64

```

Question-6:

6. Check for Categorical columns and perform encoding

Solution:

```

#Check for categorical columns & performs encoding
from sklearn.preprocessing import LabelEncoder

df['Gender'].unique()
df['Gender'].value_counts()

encoding=LabelEncoder()
df["Gender"]=encoding.fit_transform(df.iloc[:,1].values)
df

```

```

]: #Check for categorical columns & performs encoding
   from sklearn.preprocessing import LabelEncoder
   df['Gender'].unique()

```

```
]: array(['Female', 'Male'], dtype=object)
```

```

]: #Check for categorical columns & performs encoding
   df['Gender'].value_counts()

```

```

]: Male      5457
   Female    4543
   Name: Gender, dtype: int64

```

```
Out[22]:
```

| | CreditScore | Geography | Gender | Age | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|------|-------------|-----------|--------|------|--------|-----------|---------------|-----------|----------------|-----------------|--------|
| 0 | 619.0 | France | 0 | 42.0 | 2.0 | 0.00 | 1.0 | 1.0 | 1.0 | 101348.88 | 1.0 |
| 1 | 608.0 | Spain | 2 | 41.0 | 1.0 | 83807.86 | 1.0 | 0.0 | 1.0 | 112542.58 | 0.0 |
| 2 | 502.0 | France | 0 | 42.0 | 8.0 | 159660.80 | 3.0 | 1.0 | 0.0 | 113931.57 | 1.0 |
| 3 | 699.0 | France | 0 | 39.0 | 1.0 | 0.00 | 2.0 | 0.0 | 0.0 | 93826.63 | 0.0 |
| 4 | 850.0 | Spain | 2 | 43.0 | 2.0 | 125510.82 | 1.0 | 1.0 | 1.0 | 79084.10 | 0.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9995 | 771.0 | France | 0 | 39.0 | 5.0 | 0.00 | 2.0 | 1.0 | 0.0 | 96270.64 | 0.0 |
| 9996 | 516.0 | France | 0 | 35.0 | 10.0 | 57369.61 | 1.0 | 1.0 | 1.0 | 101699.77 | 0.0 |
| 9997 | 709.0 | France | 0 | 36.0 | 7.0 | 0.00 | 1.0 | 0.0 | 1.0 | 42085.58 | 1.0 |
| 9998 | 772.0 | Germany | 1 | 42.0 | 3.0 | 75075.31 | 2.0 | 1.0 | 0.0 | 92888.52 | 1.0 |
| 9999 | 792.0 | France | 0 | 28.0 | 4.0 | 130142.79 | 1.0 | 1.0 | 0.0 | 38190.78 | 0.0 |

10000 rows × 11 columns

Question-7:

7.Split the data into dependent and independent variables.

Solution:

```
#Split the data into Dependent & Independent Variables
print("-----Dependent Variables-----")
X=df.iloc[:,1:4]
print(X)
print("-----")
print("-----Independent Variables-----")
Y=df.iloc[:,4]
print(Y)
print("-----")
```

Question-8:

8. Scale the independent variables

Solution:

```
#Split the data into Dependent & Independent Variables
print("-----Dependent Variables-----")
X=df.iloc[:,1:4]
print(X)
print("-----")
print("-----Independent Variables-----")
Y=df.iloc[:,4]
print(Y)
print("-----")
```

Question-9:

9. Split the data into training and testing

Solution:

```
#Split the data into training & testing
from sklearn.model_selection import train_test_split
```

In [34]:

```
#Split the data into training & testing
x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=4, random_state=4)
x_train
x_test
y_train
y_test
Out[31]:
```

| | const | EstimatedSalary |
|------|-------|-----------------|
| 1603 | 1.0 | 23305.85 |
| 8713 | 1.0 | 41248.80 |
| 4561 | 1.0 | 143317.42 |
| 6600 | 1.0 | 174123.16 |

```
In [32]: #Split the data into training & testing
y_train
```

```
Out[32]: 2558    727.0
7642    811.0
8912    623.0
3319    430.0
6852    600.0
...
456     733.0
6017    487.0
709     686.0
8366    637.0
1146    614.0
Name: CreditScore, Length: 9996, dtype: float64
```

Out[34]:

| | const | EstimatedSalary |
|------|-------|-----------------|
| 2558 | 1.0 | 137903.54 |
| 7642 | 1.0 | 121765.00 |
| 8912 | 1.0 | 109470.34 |
| 3319 | 1.0 | 2923.61 |
| 6852 | 1.0 | 7312.25 |
| ... | ... | ... |
| 456 | 1.0 | 7666.73 |
| 6017 | 1.0 | 9085.00 |
| 709 | 1.0 | 147794.63 |
| 8366 | 1.0 | 102515.42 |
| 1146 | 1.0 | 54776.64 |

9996 rows × 2 columns